

CptS 475/575: Data Science, Fall 2020

Assignment 5: Regression and Classification

Release Date: October 16, 2020

Due Date: October 30, 2020 (11:59 pm)

General instruction: The first part of this assignment will assess your understanding of linear and logistic regression. The second part of this assignment will require you to take a set of articles from a real world newspaper and classify them based on which section they belong to. The problem is broken into four sections with points assigned to each.

Your solution will be submitted as a PDF file, which must **include your full, functional code and relevant results as stated in each part**. You are encouraged to use R Markdown to prepare your file.

- 1) This question involves the use of multiple linear regression on the **Auto** data set from the course webpage (<https://scads.eecs.wsu.edu/index.php/datasets/>). Ensure that you remove missing values from the dataframe, and that values are represented in the appropriate types.
 - a. (5%) Perform a multiple linear regression with **mpg** as the response and all other variables except **name** as the predictors. Show a printout of the result (including coefficient, error and t values for each predictor). Comment on the output:
 - i) Which predictors appear to have a statistically significant relationship to the response, and how do you determine this?
 - ii) What does the coefficient for the **displacement** variable suggest, in simple terms?
 - b. (5%) Produce diagnostic plots of the linear regression fit. Comment on any problems you see with the fit. Do the residual plots suggest any unusually large outliers? Does the leverage plot identify any observations with unusually high leverage?
 - c. (5%) Fit linear regression models with interaction effects. Do any interactions appear to be statistically significant?
- 2) This problem involves the **Boston** data set, which we saw in class. We will now try to predict per capita crime rate using the other variables in this data set. In other words, per capita crime rate is the response, and the other variables are the predictors.
 - a. (6%) For each predictor, fit a simple linear regression model to predict the response. Include the code, but not the output for all models in your solution.
 - b. (6%) In which of the models is there a statistically significant association between the predictor and the response? Considering the meaning of each variable, discuss the relationship between **crim** and **nox**, **chas**, **medv** and **dis** in particular. How do these relationships differ?
 - c. (6%) Fit a multiple regression model to predict the response using all the predictors. Describe your results. For which predictors can we reject the null hypothesis $H_0 : \beta_j = 0$?
 - d. (6%) How do your results from (a) compare to your results from (b)? Create a plot displaying the univariate regression coefficients from (a) on the x-axis, and the multiple regression coefficients from (b) on the y-axis. That is, each predictor is displayed as a single point in the plot. Its coefficient in a simple linear regression model is shown on the

x-axis, and its coefficient estimate in the multiple linear regression model is shown on the y-axis. What does this plot tell you about the various predictors?

- e. (6%) Is there evidence of non-linear association between any of the predictors and the response? To answer this question, for each predictor X , fit a model of the form

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \varepsilon$$

Hint: use the `poly()` function in R. Again, include the code, but not the output for each model in your solution, and instead describe any non-linear trends you uncover.

- 3) Suppose we collect data for a group of students in a statistics class with variables:

X_1 = hours studied,

X_2 = undergrad GPA,

X_3 = PSQI score (a sleep quality index), and

Y = receive an A.

We fit a logistic regression and produce estimated coefficient, $\beta_0 = -7$, $\beta_1 = 0.1$, $\beta_2 = 1$, $\beta_3 = -.04$.

- a. (5%) Estimate the probability that a student who studies for 32 h, has a PSQI score of 12 and has an undergrad GPA of 3.0 gets an A in the class. Show your work.
- b. (5%) How many hours would the student in part (a) need to study to have a 50 % chance of getting an A in the class? Show your work.
- c. (5%) How many hours would a student with a 3.0 GPA and a PSQI score of 3 need to study to have a 50 % chance of getting an A in the class? Show your work.
- 4) For this question, you will use a naïve Bayes model to classify newspaper articles by their section. You will be provided a set of news articles (<http://scads.eecs.wsu.edu/index.php/datasets>) collected from the Guardian (a British newspaper). The articles are cleared of major confounding factors, such as HTML tags, but it is up to you to check the articles for other problems and to prepare them for classification.
- a. Tokenization (20%)
- In order to use Naïve Bayes effectively, you will need to split your text into tokens. It is common practice when doing this to reduce your words to their stems so that conjugations produce less noise in your data. For example, the words "speak", "spoke", and "speaking" are all likely to denote a similar context, and so a stemmed tokenization will merge all of them into a single stem. R has several libraries for tokenization, stemming and text mining. Some you may want to use as a starting point are `tokenizers`, `SnowballC`, `tm` respectively, or alternatively `quanteda`, which will handle the aforementioned along with building your model in the next step. You will need to produce a document-term matrix from your stemmed tokenized data. This will have a very wide feature set (to be reduced in the following step) where each word stem is a feature, and each article has a list of values representing the number of occurrences of each stem in its body. Before representing the feature set in a non-compact storage format

(such as a plain matrix), you will want to remove any word which appears in too few documents (typically fewer than 1% of documents, but you can be more or less stringent as you see fit). You may also use a boolean for word presence/absence if you find it more effective. To demonstrate your completion of this part, you can simply select and print the text of a random article along with the **non-zero** entries of its feature vector.

b. Classification (20%)

For the final portion of this assignment, you will build and test a Naïve Bayes classifier with your data. First, you will need to use feature selection to reduce your feature set. A popular library for this is **caret**. It has many functionalities for reducing feature sets, including removing highly correlated features. You may wish to try several different methods to see which produces the best results for the following steps.

Next, you will split your data into a training set and a test set. Your training set should comprise approximately 80% of your articles, however, you may try several sizes to find which produces the best results. Whatever way you split your training and test sets, however, you should try to ensure that your six article categories are equally represented in both sets.

Next, you will build your Naïve Bayes classifier from your training data. The **e1071** package is most commonly used for this. Finally, you can use your model to predict the categories of your test data.

Once you have produced a model that generates the best predictions you can get, **print a confusion matrix** of the results to demonstrate your completion of this task. For each class, give scores for **precision** ($\text{TruePositives} / \text{TruePositives} + \text{FalsePositives}$) and **recall** ($\text{TruePositives} / \text{TruePositives} + \text{FalseNegatives}$). To do this, you may want to use the **confusionMatrix()** function.